

Presidents Game Playing Agent

Nicholas Larus-Stone

Juan Perdomo

Matt Goldberg

December 9, 2015

1 Introduction

A description of the purpose, goals, and scope of your system or empirical investigation. You should include references to papers you read on which your project and any algorithms you used are based. Include a discussion of whether you adapted a published algorithm or devised a new one, the range of problems and issues you addressed, and the relation of these problems and issues to the techniques and ideas covered in the course.

Investigating the best way to make an AI agent that plays the game Presidents. Chose this because of multiplayer games and uncertainty. Sturtevant was primary expert on this, see his thesis [here](#). Also see MCTS paper. Adapted published algorithms, with inventions based on how to deal with uncertainty (sampling and heuristic). Built a lot on ideas from two player game in course, but more complicated by expanding them.

2 Background

Research into game playing has been a main focus of members of the artificial intelligence community since the mid 20th century. As mentioned in lecture by Prof. Rush, computer scientists worked on topics such as chess AI from 1959 ¹. Although research into game playing algorithms might seem like an inefficient investment of resources by talented researchers, breakthroughs in the algorithms and techniques designed for simple games such as chess or Go have often had repercussions in other fields that have more direct impact on society. At their essence, these algorithms work to find solutions to problems in which players compete to maximize their rewards under certain constraints. In the case of Presidents, however, there is the additional complication of having imperfect information as each individual agent must operate without knowing the hands of his opponents. Developments in game playing algorithms for problems like Presidents therefore can provide insights into different real life problems such as auctions, elections, AND X in which a number of agents compete to achieve certain goals using incomplete information.

The first two algorithms discussed in this paper, Paranoid and Max-N are essentially extensions of the traditional MiniMax algorithm developed by John von Neumann in 1928 ². Since its origins, however, several optimizations have been introduced such as the use of Alpha Beta

¹Rush, Adversarial Search and Games

²Kjeldsen, John von Neumann's Conception of the Minimax Theorem: A Journey Through Different Mathematical Contexts

Pruning to reduce the size of the search space and techniques to deal with imperfect information. Moreover, research into techniques to deal with uncertainty has become a hall mark of 21st century computer science and artificial intelligence. Conversely to Paranoid and Minimax, Monte Carlo Tree Search has been developed more recently, first appearing in a paper by Remi Coulomb in 2006. Although Monte Carlo methods have their origins in the 1940s and the Manhattan Project, the creation of the MCTS algorithm for games is a 21st century discovery.

<http://www.theoremoftheday.org/Docs/Kjeldsen.pdf>, Gray

3 Related Work

During our research into game playing methods for Presidents. We found no algorithms that dealt specifically with the game of presidents, however, we did find research in related fields such as card playing AI agents and agents for games of incomplete information. Multiplayer Games, Algorithms and Approaches by Sturtevant was particularly helpful to learn about extensions of Minimax to multi agent games. Koller's work on imperfect information provided insights into adaptations of traditional game playing agents for scenarios in which there is imperfect information. Lastly, papers by Fujita and Browne on Reinforcement Learning agents and Monte Carlo Tree Search methods served as guides that led our initial development of the algorithms used.

For instance, [?].

4 Overview of Algorithms

Sampling based on cards that haven't been played

4.1 Paranoid

Modified minimax, where leaf/terminal nodes are a tuple of values representing the scores (or a heuristic estimate of those scores). Calculate intermediate values by subtracting all scores from first player score. Too deep to expand full game tree, so terminate at terminal node or after a certain depth or a certain number of nodes has been expanded (keep track of nodes expanded for pruning).

4.1.1 Pruning

Standard a-b pruning works here

4.2 Max-n

4.3 fasdfa

4.4 Monte Carlo Tree Search

A clear specification of the algorithm(s) you used and a description of the main data structures in the implementation. Include a discussion of any details of the algorithm that were not in the published paper(s) that formed the basis of your implementation. A reader should be able to reconstruct and verify your work from reading your paper.

The MCTS algorithm is divided into 4 key stages: selection, expansion, simulation, and back-propagation. In selection, a node is chosen starting from the root of the game tree by recursively selecting the most promising child node until a leaf in the tree is reached. After a node has been chosen, the game is played out using a default policy until a result is reached. Lastly, in backpropagation, the result of the playout is incorporated recursively into each node located in the path from the root to the selected child.

Algorithm 1 Pseudocode for Monte Carlo Tree Search

```

procedure MCTS(state)
  root  $\leftarrow$  mctsNode(state)
  while budget do
    nextNode  $\leftarrow$  selection(root)
    result  $\leftarrow$  simulation(nextNode)
    backpropagate(nextNode, result)
  end while
  action  $\leftarrow$  bestChild(root)
  return action
end procedure

```

In the context of Presidents, a separate tree data structure to the state formalism had to be developed in order to implement MCTS. Each node in the tree contained variables describing the number of times each node had been visited as well as the score of the node. Moreover, each node contained a list of hands for each player and the id of the player whose turn it was to play. Since the hands of the other players in the game are unknown, when instantiating the tree, hands had to be sampled for the other agents in the game based on the sizes of each player's hand and the set of cards yet to be played. Each child node corresponded to the actions the player whose turn it was could make based on his hand and the top card on the deck. Scores were defined as the finishing position of the agent in the game. Since higher scores are better, we took the inverse of the finishing rank to indicate the score.

In the selection stage, we used the UCT algorithm developed by Kocsis and Szepesvari discussed by Browne in his paper on MCTS methods. It takes into account exploration and exploitation of paths down the tree by weighing the score of each node and the number of times it and its parent had been visited. In simulation a game was played out based on the hands of each player at that node. Finally, in the backpropagation stage the normalized score (raw score divided by number of visits) was sent up the tree to the node. This is run until a predetermined budget is used up, which in this case we chose to be time. After the time ran out, the agent made a move based on the child of the action with the best score.

5 Body 2

Unsure if we want this?

Score
Approach 1
Approach 2

Table 1: Description of the results.

6 Experiments

Analysis, evaluation, and critique of the algorithm and your implementation. Include a description of the testing data you used and a discussion of examples that illustrate major features of your system. Testing is a critical part of system construction, and the scope of your testing will be an important component in our evaluation. Discuss what you learned from the implementation.

Can test on a number of different axes—time to run, nodes expanded, how it does playing against dummy agents, how it does playing against other algorithms, how it does playing against itself, number of times we sample.

6.1 Methods and Models

Wrote dummy agents to play lowest legal card-naive algorithm. Baseline to test against that for how good our agents were.

6.2 Results

For algorithm-comparison projects: a section reporting empirical comparison results preferably presented graphically.

Table 1: showing time to run on 50 games, average score

Graph 1: Paranoid (5 trials, include error bars)

Bar graph showing:

Average score for Dummy vs Dummies

Paranoid vs Dummies (sample once, 500 nodes)

Paranoid vs Dummies (sample 5 times, 500 nodes)

Paranoid vs Dummies (sample 25 times, 500 nodes)

Paranoid vs Dummies (sample 25 times, 1500 nodes)

Graph 2: Max-n

Graph 3: MCTS

6.3 Discussion

How we could improve. How important the heuristic was. RL and some of the challenges (discretizing state space)

A Program Trace

Appendix 1 A trace of the program showing how it handles key examples or some other demonstration of the program in action.

B System Description

Appendix 2 A clear description of how to use your system and how to generate the output you discussed in the write-up and the example transcript in Appendix 1. N.B.: The teaching staff must be able to run your system.

C Group Makeup

- Matt Goldberg - Wrote the state, agent, and game modules that provide the necessary formalisms and rules to encode the game of Presidents. Moreover, implemented the Max-N algorithm.
- Nicholas Larus-Stone - Implemented the repeated games function that allows a list of agents to play multiple games. Designed the Paranoid algorithm.
- Juan Perdomo- Implemented the formalisms necessary to implement the MCTS algorithm as well as implemented the algorithm itself.